

Skolemization

So far, converting wff to CNF ignored existentials

e.g. $\exists x \forall y \exists z P(x, y, z)$

Idea: names for individuals claimed to exist, called Skolem constant and function symbols

there exists an x , call it a

for each y , there is a z , call it $f(y)$

get $\forall y P(a, y, f(y))$

So replace $\forall x_1 (\dots \forall x_2 (\dots \forall x_n (\dots \exists y [\dots y \dots] \dots) \dots) \dots)$
by $\forall x_1 (\dots \forall x_2 (\dots \forall x_n (\dots [\dots f(x_1, x_2, \dots, x_n) \dots] \dots) \dots) \dots)$

f is a new function symbol that appears nowhere else

Skolemization does not preserve equivalence

e.g. $\not\models \exists x P(x) \equiv P(a)$

But it does preserve satisfiability

α is satisfiable iff α' is satisfiable (where α' is the result of Skolemization)

sufficient for resolution!

Variable dependence

Show that $\exists x \forall y R(x,y) \models \forall y \exists x R(x,y)$

show $\{\exists x \forall y R(x,y), \neg \forall y \exists x R(x,y)\}$ unsatisfiable

$$\exists x \forall y R(x,y) \implies \forall y R(a,y)$$

$$\neg \forall y \exists x R(x,y) \implies \exists y \forall x \neg R(x,y) \implies \forall x \neg R(x,b)$$

then $\{ [R(a,y)], [\neg R(x,b)] \} \rightarrow []$ with $\{x/a, y/b\}$.

Show that $\forall y \exists x R(x,y) \not\models \exists x \forall y R(x,y)$

show $\{\forall y \exists x R(x,y), \neg \exists x \forall y R(x,y)\}$ satisfiable

$$\forall y \exists x R(x,y) \implies \forall y R(f(y),y)$$

$$\neg \exists x \forall y R(x,y) \implies \forall x \exists y \neg R(x,y) \implies \forall x \neg R(x,g(x))$$

then get $\{ [R(f(y),y)], [\neg R(x,g(x))] \}$

where the two literals do not unify

Note: important to get dependence of variables correct

$R(f(y),y)$ vs. $R(a,y)$ in the above

A problem

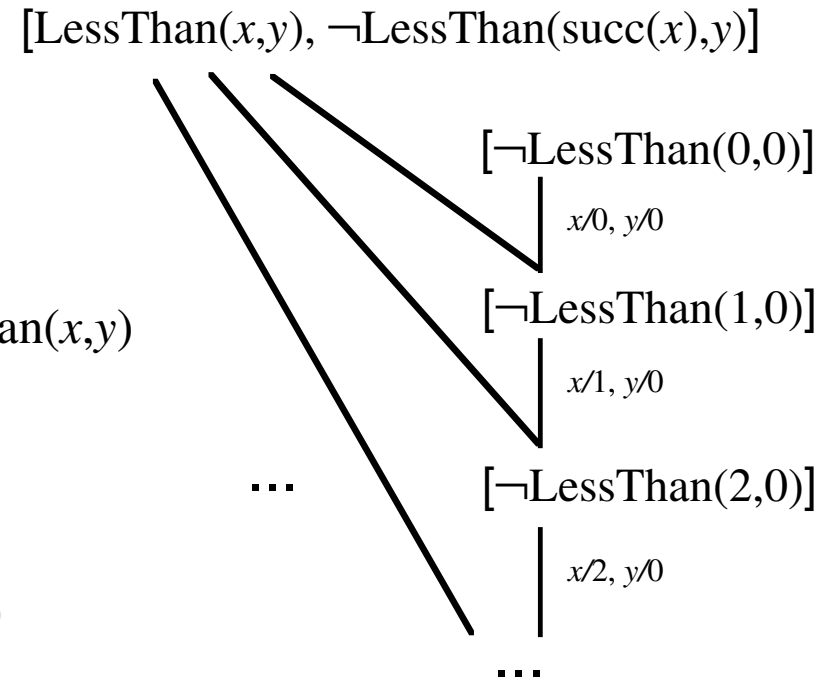
KB:

$\text{LessThan}(\text{succ}(x), y) \supset \text{LessThan}(x, y)$

Query:

$\text{LessThan}(\text{zero}, \text{zero})$

Should fail since $\text{KB} \not\models Q$



Infinite branch of resolvents

cannot use a simple depth-first
procedure to search for \square

Undecidability

Is there a way to detect when this happens?

No! FOL is very powerful

can be used as a full programming language

just as there is no way to detect in general when
a program is looping

There can be no procedure that does this:

Proc[*Clauses*] =


 If *Clauses* are unsatisfiable
 then return YES
 else return NO

However: Resolution is complete

some branch will contain [], for unsatisfiable clauses

So breadth-first search guaranteed to find []

search may not terminate on satisfiable clauses



Overly specific unifiers

In general, no way to guarantee efficiency, or even termination

later: put control into users' hands

One thing that can be done:

reduce redundancy in search, by keeping search as general as possible

Example

$..., P(g(x), f(x), z)] \quad [\neg P(y, f(w), a), ...$

unified by

$\theta_1 = \{x/b, y/g(b), z/a, w/b\}$ gives $P(g(b), f(b), a)$

and by

$\theta_2 = \{x/f(z), y/g(f(z)), z/a, w/f(z)\}$ gives $P(g(f(z)), f(f(z)), a)$.

Might not be able to derive the empty clause from clauses having overly specific substitutions

wastes time in search!

Most general unifiers

θ is a most general unifier (MGU) of literals ρ_1 and ρ_2 iff

1. θ unifies ρ_1 and ρ_2
2. for any other unifier θ' , there is a another substitution θ^* such that $\theta' = \theta\theta^*$

Note: composition $\theta\theta^*$ requires applying θ^* to terms in θ

for previous example, an MGU is

$$\theta = \{x/w, y/g(w), z/a\}$$

for which

$$\theta_1 = \theta\{w/b\}$$

$$\theta_2 = \theta\{w/f(z)\}$$

Theorem: Can limit search to most general unifiers only without loss of completeness (with certain caveats)

Computing MGUs

Computing an MGU, given a set of literals $\{\rho_i\}$

usually only have two literals

1. Start with $\theta := \{\}$.
2. If all the $\rho_i\theta$ are identical, then done;
otherwise, get disagreement set, DS
e.g. $P(a, f(a, g(z)), \dots) \quad P(a, f(a, u), \dots)$
disagreement set, $DS = \{u, g(z)\}$
3. Find a variable $v \in DS$, and a term $t \in DS$ not containing v .
If not, fail.
4. $\theta := \theta\{v/t\}$
5. Go to 2

Note: there is a better *linear* algorithm

Herbrand Theorem

Some 1st-order cases can be handled by converting them to a propositional form

Given a set of clauses S

- the Herbrand universe of S is the set of all terms formed using only the function symbols in S (at least one)
e.g., if S uses (unary) f , and c, d , $U = \{c, d, f(c), f(d), f(f(c)), f(f(d)), f(f(f(c))), \dots\}$
- the Herbrand base of S is the set of all $c\theta$ such that $c \in S$ and θ replaces the variables in c by terms from the Herbrand universe

Theorem: S is satisfiable iff Herbrand base is

(applies to Horn clauses also)

Herbrand base has no variables, and so is essentially *propositional*, though usually infinite

- finite, when Herbrand universe is finite
can use propositional methods (guaranteed to terminate)
- sometimes other “type” restrictions can be used to keep the Herbrand base finite
include $f(t)$ only if t is the correct type

Resolution is difficult!

First-order resolution is not guaranteed to terminate.

What can be said about the propositional case?

Shown by Haken in 1985 that there are unsatisfiable clauses $\{c_1, c_2, \dots, c_n\}$ such that the *shortest* derivation of \square contains on the order of 2^n clauses

Even if we could always find a derivation immediately, the most clever search procedure will still require *exponential* time on some problems

Problem just with resolution?

Probably not.

Determining if a set of clauses is satisfiable was shown by Cook in 1972 to be NP-complete

No easier than an extremely large variety of computational tasks

Roughly: any search task where what is searched for can be verified in polynomial time can be recast as a satisfiability problem

- » satisfiability
- » does graph of cities allow for a full tour of size $\leq k$ miles?
- » can N queens be put on an $N \times N$ chessboard all safely? and many, many more....

Satisfiability is believed by most people to be unsolvable in polynomial time

SAT solvers

In the propositional case, procedures have been proposed for determining the satisfiability of a set of clauses that appear to work much better in practice than Resolution.

The most popular is called DP (or DPLL) based on ideas by Davis, Putnam, Loveland and Logemann. (See book for details.)

These procedures are called SAT solvers as they are mostly used to find a satisfying interpretation for clauses that are satisfiable.

related to constraint satisfaction programs (CSP)

Typically they have the property that if they *fail* to find a satisfying interpretation, a Resolution derivation of [] can be reconstructed from a trace of their execution.

so worst-case exponential behaviour, via Haken's theorem!

One interesting counter-example to this is the procedure GSAT, which has different limitations. (Again, see the book.)

Implications for KR

Problem: want to produce entailments of KB as needed for immediate action

full theorem-proving may be too difficult for KR!

need to consider other options ...

- giving control to user e.g. procedural representations (later)
- less expressive languages e.g. Horn clauses (and a major theme later)

In some applications, it is reasonable to wait

e.g. mathematical theorem proving, where we care about specific formulas

Best to hope for in general: reduce redundancy

main example: MGU, as before

but many other strategies (as we will see)

ATP: automated theorem proving

- area of AI that studies strategies for automatically proving difficult theorems
- main application: mathematics, but relevance also to KR

Strategies

1. Clause elimination

- pure clause

contains literal p such that p does not appear in any other clause
clause cannot lead to \square

- tautology

clause with a literal and its negation
any path to \square can bypass tautology

- subsumed clause

a clause such that one with a subset of its literals is already present
path to \square need only pass through short clause
can be generalized to allow substitutions

2. Ordering strategies

many possible ways to order search, but best and simplest is

- unit preference

prefer to resolve unit clauses first

Why? Given unit clause and another clause, resolvent is a smaller one $\Rightarrow \square$

Strategies 2

3. Set of support

KB is usually satisfiable, so not very useful to resolve among clauses with only ancestors in KB

contradiction arises from interaction with $\neg Q$

always resolve with at least one clause that has an ancestor in $\neg Q$

preserves completeness (sometimes)

4. Connection graph

pre-compute all possible unifications

build a graph with edges between any two unifiable literals of opposite polarity

label edge with MGU

Resolution procedure:

repeatedly: select link
 compute resolvent
 inherit links from parents after substitution

Resolution as search: find sequence of links L_1, L_2, \dots producing []

Strategies 3

5. Special treatment for equality

instead of using axioms for =

reflexivity, symmetry, transitivity, substitution of equals for equals

use new inference rule: paramodulation

from $\{(t=s)\} \cup C_1$ and $\{P(\dots t' \dots)\} \cup C_2$

where $t\theta = t'\theta$

infer $\{P(\dots s \dots)\}\theta \cup C_1\theta \cup C_2\theta$.

collapses many resolution steps into one

see also: theory resolution (later)

6. Sorted logic

terms get sorts:

x : Male mother:[Person \rightarrow Female]

keep taxonomy of sorts

only unify $P(s)$ with $P(t)$ when sorts are compatible

assumes only “meaningful” paths will lead to []

Finally...

7. Directional connectives

given $[\neg p, q]$, can interpret as either

from p , infer q (forward)

to prove q , prove p (backward)

procedural reading of \supset

In 1st case: would only resolve $[\neg p, q]$ with $[p, \dots]$ producing $[q, \dots]$

In 2nd case: would only resolve $[\neg p, q]$ with $[\neg q, \dots]$ producing $[\neg p, \dots]$

Intended application:

forward: Battleship(x) \supset Gray(x)

do not want to try to prove something is gray
by trying to prove that it is a battleship

backward: Person(x) \supset Has(x , spleen)

do not want to conclude the spleen property for
each individual inferred to be a person

This is the starting point for the procedural representations (later)